

(1)

IAN Goodfellow - Deep learning

14 = Autoencoders

15 = Representation learning

20.10.8/4 = Deep gen models -
VAEs & GANs

Autoencoders

An α Autoencoder is a nn that is trained to attempt to copy its input to its output

Internally, a hidden layer h describes a code used to represent the input
 \uparrow latent

Network of two parts:

(1) An encoder $h = f(x)$

(2) A decoder $r = g(h)$



The ability to copy x exactly is not useful

Autoencoders are designed to be unable to learn to copy perfectly

copy to resemble train data

(2)

traditionally (historically) Autoencoders were used for dim reduction or feature learning

Recent theory connecting Autoencoder & latent variable models

Autoencoders may be thought as a special case of feedforward networks

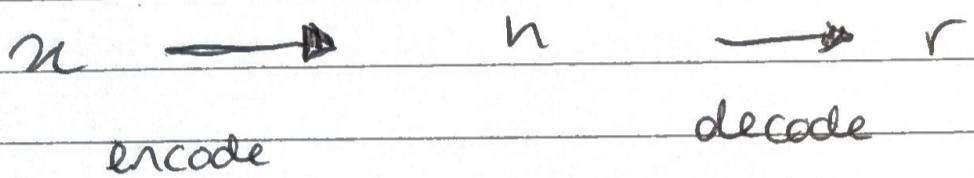
- ↳ & may be trained w/ all same techs
 - ↳ minibatch gradient descent, Backprop
-

14. 1 undercomplete Autoencoders
 14. 2 Regularized Autoencoders
 14. 2. 1 Sparse Autoencoders
 14. 2. 2 Denoising Autoencoders
 14. 2. 3 regularize by penalizing derivatives
 14. 3 rep power, layer size & depth
 14. 4 Stochastic encoders & decoders
 14. 5 Denoising Autoencoders
 14. 6 learning manifolds w/ Autoencoders
 14. 7 Contractive autoencoders
 14. 8 Predictive sparse decomposition
 14. 9 Applications of Autoencoders
-

(3)

14.1 undercomplete Autoencoders

whilst achieving a copy from input to output may sound trivial we hope that the intermediary task of h will take on some useful properties



one way to make it useful is to constrain h to be smaller than n

↳ this is why it is called undercomplete

undercomplete forces the Autoencoder to capture the most salient features

learn progress \rightarrow min loss function

$$L(n, g(f(x)))$$

loss func penalizes $g(f(x))$ for being dissimilar to x

encoder & Decoder can't be allow too much capacity otherwise the learning will just be an identity w/ no features learnt

14.2 Regularized Autoencoders

- undercomplete
 - code dimension less than input
 - learn salient features
 - will not learn anything if encoder + Decoder = too much capacity

Same issue occurs if hidden code has same dimensions as input

- or hidden has more dims than input = overcomplete

Ideally, train Autoencoder successfully by choosing code dims & capacity based on complexity of data to be modelled

Reg AutoEncoders provide ability

rather than limiting the capacity by keeping the en/de shallow

- Reg Autoencoders use a loss function that encourages the model to have other properties besides the ability to copy

A regularized autoencoder can still learn something useful about the data even if overcomplete

~~Notes~~
VAEs (variational autoencoders) are regularized

► 14.2.1 Sparse Autoencoders

An autoencoder that includes a sparsity penalty $\mathcal{L}(h)$ on the code/hidden layer

$$\mathcal{L}(n, g(f(n))) + \mathcal{L}(h)$$

$g()$ Decoder output

Sparse encoders learn features to then be applied for another task

► Classyaction fe

must respond to unique statistical features rather than becoming an identity func

thus training to perform the copy task w/ a sparsity penalty can yield a model that has learned

take a feedforward network whose task is to copy the input

► add a sparse reg term

Regularized Autoencoders are not Bayesian as the reg depends on the data

C

rather than think of sparsity penalty as a reg for the copy task

think of entire sparse Autoencoder as approximating Max likelihood of train a generative model that has latent variables

14.3 Representational Power, layer size & Depth

often Autoencoders =

single layer encode + single layer decode

However, this is not a requirement

using deep encoders & decoders has many advantages

Note: ~~Feedforward NN's~~ benefit from Depth

↳ Autoencoders are feedforwards
↳ encoders & decoders are FFS

↳ All components benefit from Depth

Advantage of depth; universal Approx theorem guarantees that a feedforw network w/ atleast one hidden layer can represent an approx of any func

(7)

to an arbitrary degree of accuracy
provided enough hidden units

Arbitrary here means any

↳ there exists a num of hid units
to achieve the selected accuracy

if an encoder is shallow - no hidden -
cannot achieve arb constraints

A deep encoder can

Depth can reduce ~~cost~~ comp cost

↳ requires less weights and bias

↳ layers build up features

↳ low → med → high

↳ shallow needs many units to
achieve this

↳ this costs more weights & bias
= comp cost ↑

→ layers can be seen as modular
& reuse-able

14.4 Stochastic Encoders & Decoders

Autoencoders are just feedforward networks

- ↳ same loss funcs can be used
- ↳ same output unit types

In trad feednet, $p(y|x)$.

↳ y is the target given x

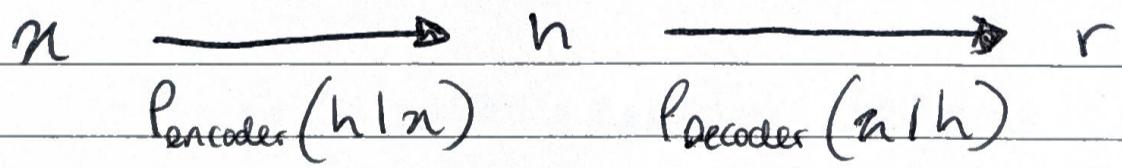
↳ y = vector of target such as class labels

In an autoencoder, x (input) is also a target for the Decoder $\xrightarrow{h} y(x)$

Encoder ($x \mapsto h$)

train, min $-\log p_{\text{decoder}}(y|h)$

Stochastic Architecture →



both encoder & decoder are not simple functions

↳ instead involve some noise injections

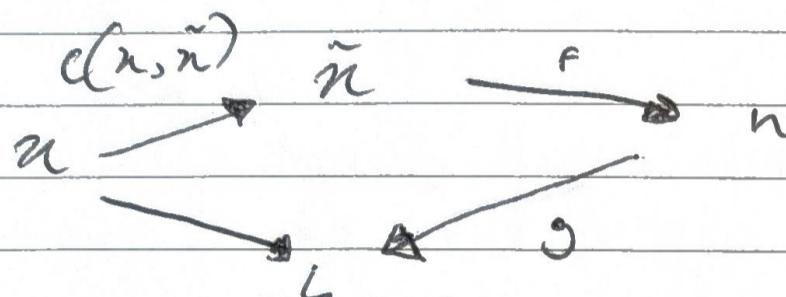
↳ thus, output can be seen as being sampled from a distribution

Any latent variable model defines a stochastic encoder/decoder

(9)

14.5 Denoising Autoencoders DAE

An A DAE receives a corrupted data point as its input \tilde{n} is trained to predict the original, uncorrupted data point n .



corruption process $C(\tilde{z}|n)$

dates back to LeCun (1987)

Denoising Autoencoders are just MLEPs trained to denoise

"Denoising Autoencoders" as a specific type of model

- ↳ not just merely denoise input
- ↳ but learn good internal representations

goal is to allow the creation of a high capacity encoder

- ↳ but avoid identity function

14.6 Manifolds

~~Very complex & disordered~~

the manifold hypothesis says

"Many high dimensional datasets
that occur in the real world
actually lay upon some low
dimensional manifold within
that high dimensional space"

2-D examples of a manifold

↳ line, circle

3-D manifolds

↳ plane, sphere

0-D = Dot

goal of A NN is to stretch & map
these & disentangle these manifolds

such that we can finally separate
~~these~~ them using a hyperplane
in the final layer

14.7 Contractive Autoencoders

14.8 Predictive sparse decomposition

Hybrid of Sparse & Parametric

Parametric encoder = trained to predict
the output of iterative inference

14.9 Applications of Autoencoders

Autoencoders have been success applied to
dimension reduction & information retrieval
tasks

Dim red = one of first applications of ~~rep~~
representation learning & DL

↳ early study motivs for Autoencoders

Why is dimension reduction good?

- ↳ lower dims improve performance on many tasks
- ↳ less memory & runtime

Many forms of Dim reduction place
semantically related example near
each other

the hints provided by mapping to lower
dimensions aid generalization