

Generative Adversarial Nets

framework for estimating gen models via an adversarial process

Simil train two models

- G : A gen model - Capture Data dist
- D : A discrimin model - est prob that same came from train Data

training for G is to make max prob of D making a mistake

framework corresponds to a minimax two player game

A unique solution exists where G recovers the training data distribution

& D equal to $\frac{1}{2}$ everywhere

If $G \approx D = MLPs$ then can be trained w/ back-prop

introduction

historically, the best DC successes involved Discriminative models

- ↳ usually mapping high dim inputs to a class label

(2)

Mostly based on Backprop, dropout, & piecewise linear units

Deep generative models less successful
due to difficult of approx many intractable probabilistic computations from max likelihood estimators

Adversarial nets framework

gen model is pitted against an adversary

Adversary = Discrimin that learns to determine whether a sample is from model dist or data dist

both models improve in the training

gen & disc training together

in feedforward disc = discarded & only generative is used

no approximate or markov chains are needed

2 Related Works

(3)

↳

3 Adversarial Nets

most simple to do when D & G are MCP's

to learn the generators dist $p_g(\cdot)$ over \mathcal{X}

we define prior input noise variables $p_z(z)$
(just produces a random noise)

then map to data space $G(z, \theta_g)$

G = differentiable func, rep by MCP w/ Params
(Params = all weights & biases)

then define a second MCP: $D(x, \theta_d)$
→ this outputs a scalar

$D(x) = \text{Prob that } x \text{ came from data}$
rather than p_g

train D to maximise prob of assigning
the correct label to both train examples
& samples from the generator

Simultaneously train G to minimize: $\log(1 - D(g(x)))$

D & G play two player mini-max game

w/ value func $v(G, D)$ $\min_G \max_D V(D, G)$

Sometimes, early in train when G is poor there is not sufficient gradient to train G

In this case $\log(1 - D(g(z)))$ saturates

rather than $\min \log(1 - D(g(z)))$

can instead max $\log D(G(z))$

provides stronger gradients in early learning

4 theoretical Results

the Generator implicitly defines a prob distribution P_g as the dist of the samples $G(z)$ when $z \sim P_z$

5 Experiments

generators uses a mix of rectifier & sigmoid activation functions

Disc used maxout activations

Dropout used for Disc

(5)

has a challenge table for Autocode vs Adversial modes

Autoencoders = Recon vs gen
 Adversary = Disc vs gen

Inference for Autoencodes = MCMC

Advantages & Disadvantages

① Disadvantage = no explicit representation of $P_g(x)$

Provides no formula or func that tells you the prob dens of any given point x

Gans learn through implicit process

Learn mapping $g(z; \theta_g)$ from noise $P_z(z)$

gen is trained to fool disc, not explain the prob density funcs of outputs

to gen must sample $P_z(z)$ and pass to get

to no pass on arbitrary x

Makes it v hard to compare GANs

(6)

tends to rely on the qualitative evals

Other ~~VAEs~~ like gen models (like VAEs)
use log-likelihood metrics

GANs don't have an Encoder -
no way to get back to input

GANs are good @ sample from complement

Poor @ density estimation of dist