

ML Neural Nets - Supervise 3

NN, Perceptron, Multi-layer perceptron

Perceptron Algorithm

Based on gradient Descent

1. Init w , i.e. set it to some value

2. Given data instance x_n compute output
 $\hat{y}_n = \text{Sign}(x_n w)$ where $\text{Sign}()$ is an activation function

3. Update w as
$$w_{\text{new}} = w_{\text{old}} - \lambda \frac{dL_{\text{Hinge}}}{dw}$$

gradient Descent loss
Differentiate the hinge loss

$$= w_{\text{old}} - \lambda y_n x_n \mathbb{I}[\hat{y}_n \neq y_n]$$

where $\lambda = 1$

$$\mathbb{I}[\hat{y}_n \neq y_n] = 1 \text{ if } \hat{y}_n \neq y_n$$

$$\mathbb{I}[\hat{y}_n \neq y_n] = 0 \text{ if } \hat{y}_n = y_n$$

4. Repeat until convergence

(2)

Stochastic Gradient Descent (SGD)

Approximate the ~~pos~~ gradient using just one data instance

before:

$$\frac{dL(w)}{dw} = \frac{1}{N} \sum_{n=1}^N \frac{dl(f_w(x_n), y_n)}{dw}$$

With SGD

$$\frac{dL(w)}{dw} \approx \frac{dl(f_w(x_t), y_t)}{dw}$$

$$\Rightarrow w^{(t+1)} = w^{(t)} - \lambda \frac{dl(f_w(x_t), y_t)}{dw}$$

where t = current iteration

$(x_t, y_t) = (x_n, y_n)$ @ current iteration

(3)

Mini-batch

The method actually used in ML fits somewhere in between called mini-batching

MB updates w/ approximate descent of B data instances

$\{(x_b, y_b), \forall b, 1 \leq b \leq B\}$, where $B \ll N$

$$\Rightarrow w^{(\tau+1)} = w^{(\tau)} - \eta \frac{1}{B} \sum_{b=1}^B \frac{dl(f_w(x_b), y_b)}{dw}$$

~~update~~ update the model using the loss gradient from this batch

Perceptron Convergence Theorem (4)

if a solution exists, i.e. linearly sep

the perc algo will find it in a finite number of iterations

~~limitation of Perceptron is linear separability~~

Summary of Perceptron

- ① it is a single layer NN
- ② Only works for binary classification
- ③ Perceptron will fail for non-linearly separable classes
- ④ Optimization of model params based on gradient Descent