

# Applied Natural Language Processing

Dr Jeff Mitchell, University of Sussex

Autumn 2025

# Document similarity and clustering

## Previously

- Binary classification scenarios e.g., sentiment and relevance
- Wordlist classifiers
- Machine learning approaches
  - Automatically deriving wordlists
  - Naïve Bayes
- Evaluation
  - accuracy and error rate
  - the confusion matrix
  - precision, recall and F1-score

## This time

- Document Similarity
  - Vector Space Model
  - TF-IDF
  - Cosine similarity
  - Beyond words as dimensions
- Clustering
  - k-means

# Document similarity

Part 2

# Document similarity scenarios

- Navigating a document collection
  - Given a large document collection (e.g., the web)
  - Find documents which are similar to ones we already know about
- Clustering a document collection
  - Given a large document collection
  - Find natural groupings of documents which seem to be about similar topics
- Searching a document collection
  - Given a large document collection
  - Find documents which are relevant to a search query

# Similarity vs classification

- Related problem to topic classification but not identical

## Classification


- Fixed number of classes
- Classes known in advance

## Similarity

- May not know the classes or even the number of classes of interest
- Documents in different classes may actually be similar
- Similarity can be used for classification by assigning the class of the most similar document(s) in a training sample.

# Vector space model for documents

- Identify a fixed set of topic terms
  - Maybe 100s of thousands of topic terms
  - Maybe all of the content words or lemmas in the vocabulary
- One dimension in space for each term
  - This is a really high-dimensional space
  - We will normally only draw 2 (or 3) but really there are 100s of thousands of dimensions
- Every document is associated with a point (or vector) in this space



# Measuring value of a term

- Weight or value of a term should reflect importance in document
- Term Frequency
  - More frequent terms may be more important
  - But high frequency words are not always discriminating
    - E.g., stop-words
    - The term “country” in a collection of documents about travel
- Document Frequency
  - Terms which occur in less documents may be more important
  - The term “Kenya” in 3 documents in a collection of documents about travel

# Term Frequency Inverse Document Frequency (TFIDF)

- Term frequency is the number of occurrences of a term in a document:

$$tf(d, t)$$

- Document frequency is the number of documents in which a term occurs:

$$df(t)$$

- If there are  $N$  documents in total, inverse document frequency is given by:

$$idf(t) = \log\left(N/df(t)\right)$$

- Term frequency inverse document frequency:


$$tfidf(d, t) = tf(d, t) \times idf(t)$$



# Variants on TFIDF


- Many variations on the original TFIDF measure e.g.,
  - Boolean “frequencies”:  $tf(d, t) = 1$  if  $t$  appears in  $d$  and 0 otherwise
  - term frequency adjusted for document length:  $tf(d, t) = \frac{freq(d, t)}{\sum_{t' \in d} freq(d, t')}$
  - logarithmically scaled term frequency:  $tf(d, t) = \log(1 + freq(d, t))$
- why do you think each of the above adjustments might be useful?

# Vector similarity



- Similar items are close together in the vector space
- Dissimilar items are far apart
- How do we measure distance or closeness in a vector space?
  - Euclidean distance
  - Cosine similarity

# Cosine similarity



- The more similar two documents are, the smaller the angle  $\theta$  between their vectors will be.
- Recall from Maths:
  - $\cos(0) = 1$
  - $\cos(90) = 0$
- So:

$$\begin{aligned} \text{sim}(d_1, d_2) &= \cos(\theta) \\ &= \frac{\underline{a} \cdot \underline{b}}{\sqrt{\underline{a} \cdot \underline{a} \times \underline{b} \cdot \underline{b}}} \end{aligned}$$

dot product

Where:

$$\underline{a} \cdot \underline{b} = \sum_i^m a_i b_i$$

m=number of dimensions

# Cosine similarity intuitions

- Dot product of two vectors is high when they are in the same direction
- For example:

$$\underline{a} = \begin{pmatrix} 5 \\ 0 \\ 1 \end{pmatrix}$$

$$\underline{b} = \begin{pmatrix} 10 \\ 1 \\ 0 \end{pmatrix}$$

$$\underline{c} = \begin{pmatrix} 1 \\ 10 \\ 1 \end{pmatrix}$$

$$\underline{a} \cdot \underline{b} = 5 \times 10 + 0 \times 1 + 1 \times 0 = 50$$

$$\underline{b} \cdot \underline{c} = ?$$

$$\underline{a} \cdot \underline{c} = 5 \times 1 + 0 \times 10 + 1 \times 1 = 6$$

# Cosine Similarity denominator

- The denominator of the cosine calculation normalizes for the lengths of the vectors, so that we get an answer between 0 and 1

$$\underline{a} = \begin{pmatrix} 5 \\ 0 \\ 1 \end{pmatrix} \quad \underline{b} = \begin{pmatrix} 10 \\ 1 \\ 0 \end{pmatrix} \quad \underline{c} = \begin{pmatrix} 1 \\ 10 \\ 1 \end{pmatrix}$$

$$\underline{a} \cdot \underline{b} = 50$$

$$\underline{a} \cdot \underline{a} = 5 \times 5 + 1 \times 1 = 26$$

$$\cos(\underline{a}, \underline{b}) = \frac{50}{\sqrt{26 \times 101}}$$

$$\underline{a} \cdot \underline{c} = 6$$

$$\underline{b} \cdot \underline{b} = 10 \times 10 + 1 \times 1 = 101$$

$$\underline{b} \cdot \underline{c} = ?$$

$$\underline{c} \cdot \underline{c} = 1 \times 1 + 10 \times 10 + 1 \times 1 = 102$$

$$= 0.976$$

# Using cosine

Measure similarity of

- Two documents
- Document and a query
- Two queries
- Two terms
  - where terms have vectors with documents as dimensions
    - Latent Semantic Analysis
  - where terms have vectors with co-occurring terms as dimensions
    - Distributional semantics

# Beyond words

# Beyond words as dimensions

- So far have assumed that topic terms are **unigrams** (single words)
- Often a phrase or multiword term is characteristic of a topic
  - **n-grams**
  - e.g., “hedge fund”, “black hole”, “surface to air missile”
  - Individual words in phrase may be ambiguous or too general
- Same TFIDF weighting can be applied to phrases



# Phrases as topic terms

- Lots of unigrams. Maybe 100 000
- Even more bigrams. Maybe  $100\,000^2$
- Some useful phrases even longer. How many possible 5 word phrases are there in the English language?
- How do we find useful ones?

# Collocations

- Collocations are n-grams which occur together more often than by chance
- Consider an observed bigram “*black hole*”
- How often does the bigram occur? How often do the individual unigrams “*black*” and “*hole*” occur?
- How frequently might we expect “*black hole*” to occur if words occurred independently of each other?
- Collocations are n-grams where the **observed joint probability** is greater than the **expected joint probability** for independent events

# Point-wise Mutual Information (PMI)

$$PMI(w_1, w_2) = \log \left( \frac{P(w_1, w_2)}{P(w_1) \cdot P(w_2)} \right)$$

observed joint probability

expected joint probability  
(assuming independence)

- PMI tells us how surprising it is that a phrase has occurred as frequently as it does
- If the **observed joint probability** > **expected joint probability**
  - Ratio greater than 1
  - PMI is positive
- If the **observed joint probability** < **expected joint probability**
  - Ratio less than 1
  - PMI is negative

**COLLOCATION**

# Recap

# TFIDF Questions


	Doc 1	Doc 2	Doc 3	Doc 4
word 1	10	15	8	12
word 2	20	0	0	0
word 3	5	7	0	0
word 4	0	3	7	8

1. What is the tfidf value of word 3 in Doc 2?
2. What kind of word is word 1 likely to be? Give an example.
3. What kind of word is word 2 likely to be? Give an example

# Clustering

# Clustering

**Clustering** (or cluster analysis) is the task of grouping objects in such a way that objects in the same group (called a **cluster**) are **more similar** to each other than to those in other clusters.

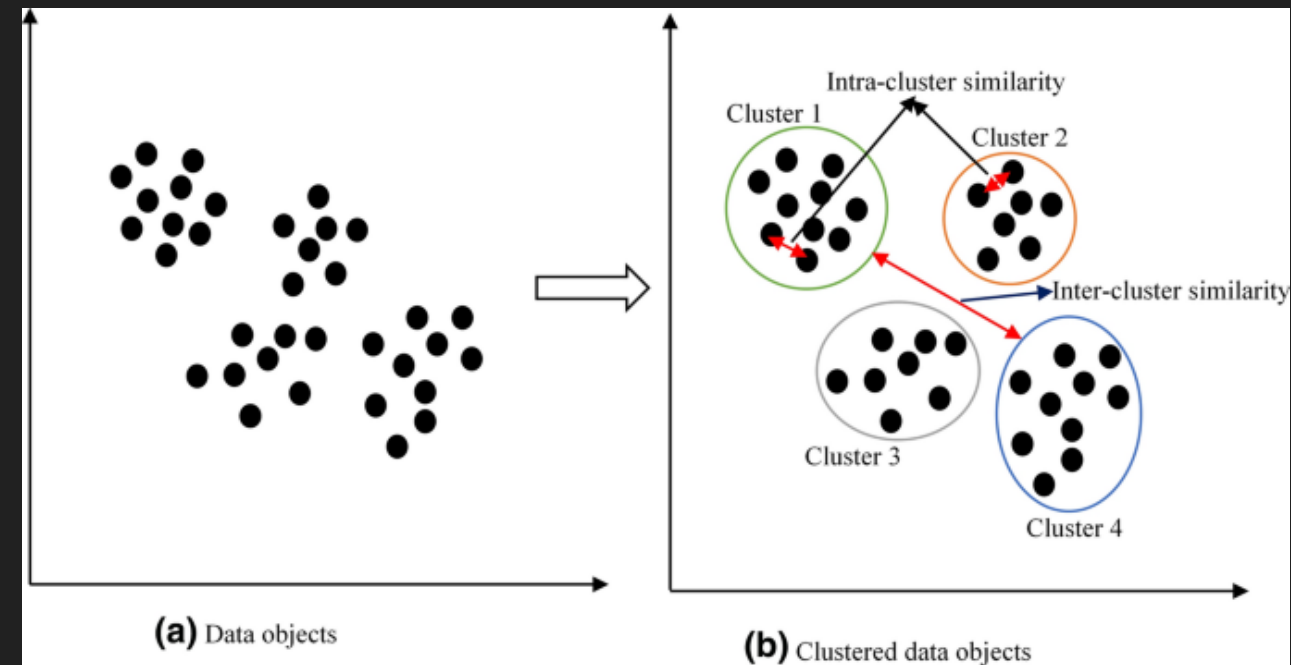


In general, there is no one “correct” clustering. For example, a different algorithm may assign the red point closest to the grey cluster to the grey cluster.

Clustering algorithms all affected by choice of similarity / distance measure and scaling of dimensions

# Clustering


- Clusters are groups of objects which are mutually similar to each other
  - => intra-cluster similarity is high
- and are dissimilar to objects outside of the cluster
  - => inter-cluster similarity is low
- objects can be documents in vector space
- how do we find the clusters?





# K-means

- Simple partitioning approach based on the idea of centroids or prototypes




The centroid (middle) of any group of points in a Euclidean space can be found by taking the mean on every dimension.

$$c_{d_j} = \frac{\sum_{i=0}^n p_{d_j,i}}{n}$$


# K-means

1. Select  $K$  points as initial centroids
2. while centroids changing:
  1. Form  $K$  clusters by assigning each point to its nearest centroid
  2. Recompute centroid of the cluster




# K-means

1. Select  $K$  points as initial centroids
2. while centroids changing:
  1. Form  $K$  clusters by assigning each point to its nearest centroid
  2. Recompute centroid of the cluster




# K-means

1. Select  $K$  points as initial centroids
2. while centroids changing:
  1. Form  $K$  clusters by assigning each point to its nearest centroid
  2. Recompute centroid of the cluster



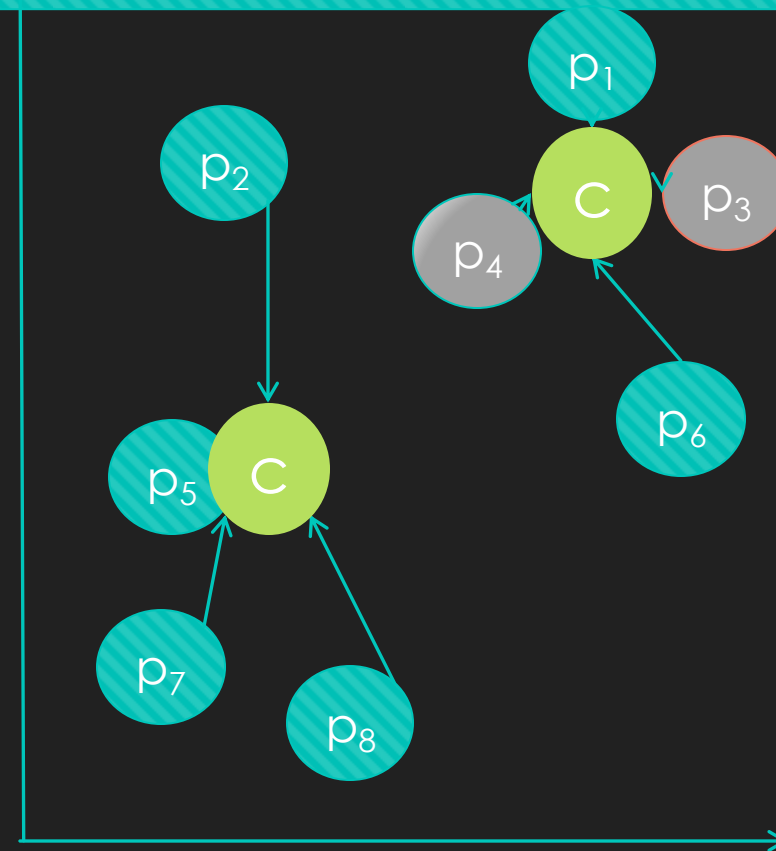
# K-means

1. Select  $K$  points as initial centroids
2. while centroids changing:
  1. Form  $K$  clusters by assigning each point to its nearest centroid
  2. Recompute centroid of the cluster



# K-means

1. Select  $K$  points as initial centroids
2. while centroids changing:
  1. Form  $K$  clusters by assigning each point to its nearest centroid
  2. Recompute centroid of the cluster



# Disadvantages of K-Means

- Need to know the number of clusters in advance
- Each point is only assigned to a single cluster (hard clustering)
- Flat structure (no clusters within clusters)
- May converge on local minimum i.e., suboptimal clustering (this can be overcome to some extent by repeated random initialisations)

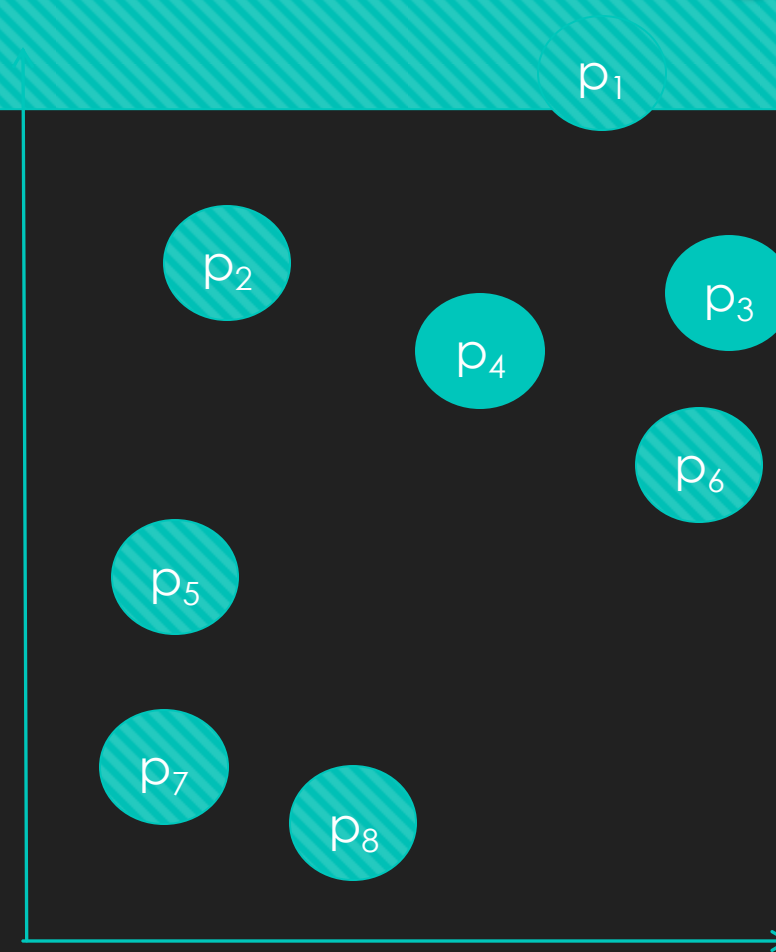
# Agglomerative hierarchical clustering

- is an agglomerative technique which builds up clusters by repeatedly merging the closest pair of clusters
- Do not need to know the number of clusters in advance
- Hierarchical (clusters have internal structure)



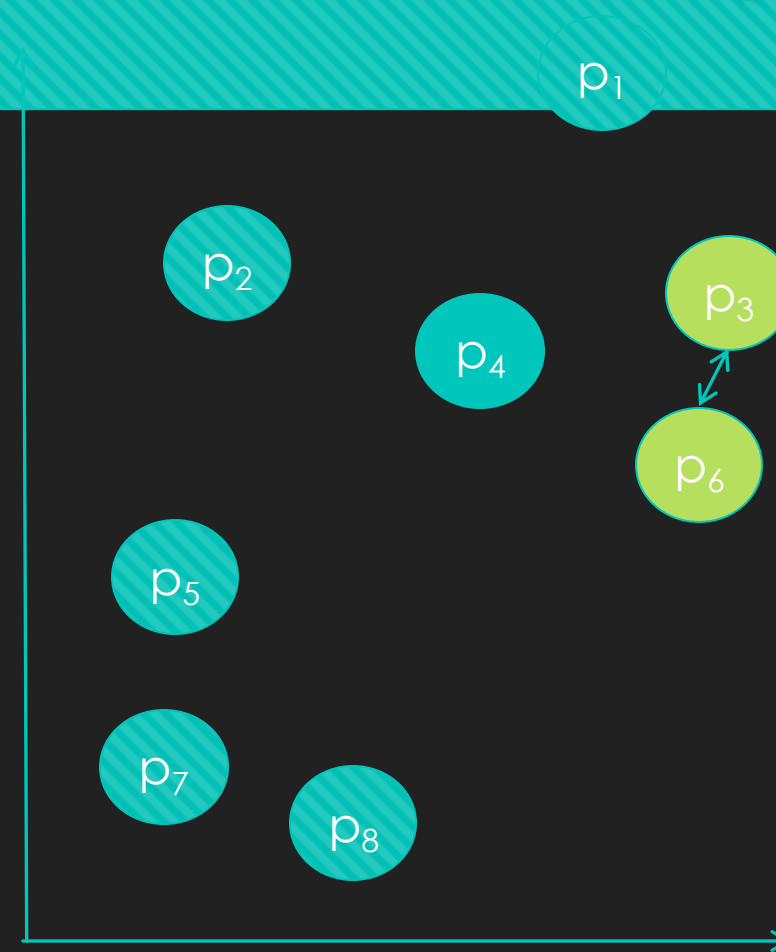
# Agglomerative hierarchical clustering

1. Initialise  $n$  clusters as the  $n$  data points
2. Find closest pair  $\langle p_i, p_j \rangle$  of clusters with distance  $d$
3. while  $d < \text{threshold}$ :
  1. Merge clusters  $\langle p_i, p_j \rangle$
  2. Find closest pair  $\langle p_i, p_j \rangle$  with distance  $d$



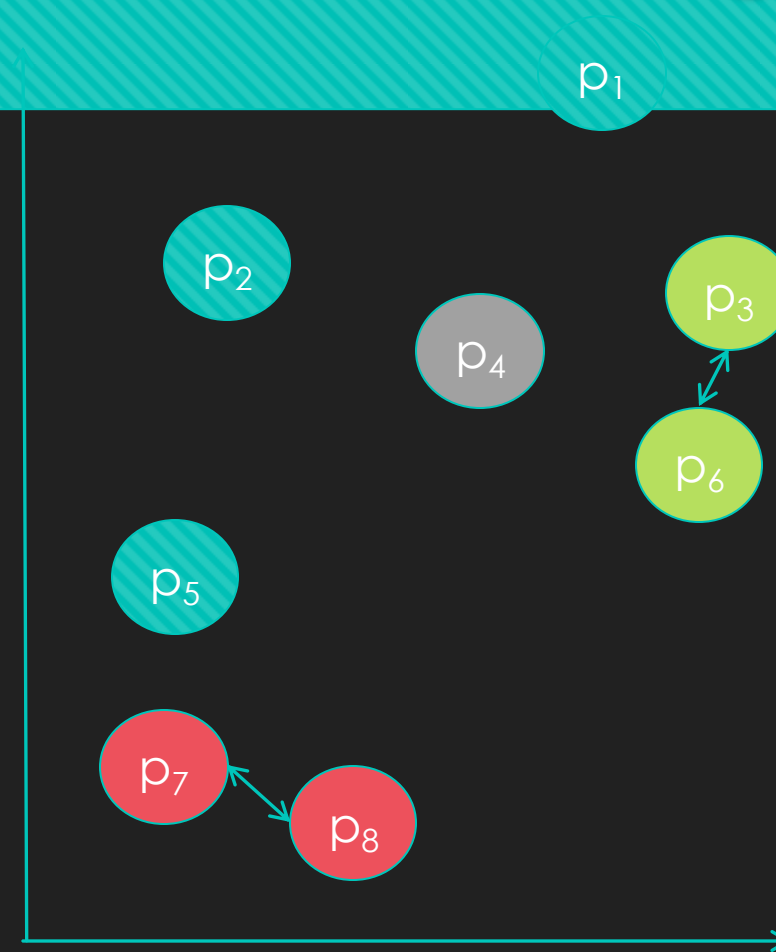
# Agglomerative hierarchical clustering

1. Initialise  $n$  clusters as the  $n$  data points
2. Find closest pair  $\langle p_i, p_j \rangle$  of clusters with distance  $d$
3. while  $d < \text{threshold}$ :
  1. Merge clusters  $\langle p_i, p_j \rangle$
  2. Find closest pair  $\langle p_i, p_j \rangle$  with distance  $d$



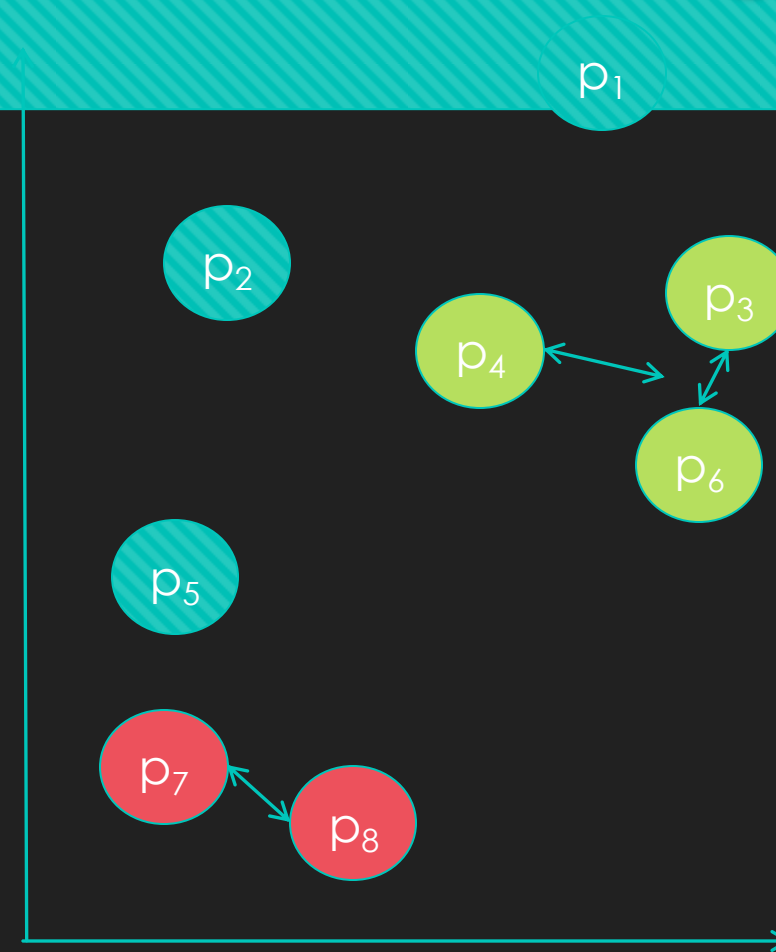
# Agglomerative hierarchical clustering

1. Initialise  $n$  clusters as the  $n$  data points
2. Find closest pair  $\langle p_i, p_j \rangle$  of clusters with distance  $d$
3. while  $d < \text{threshold}$ :
  1. Merge clusters  $\langle p_i, p_j \rangle$
  2. Find closest pair  $\langle p_i, p_j \rangle$  with distance  $d$



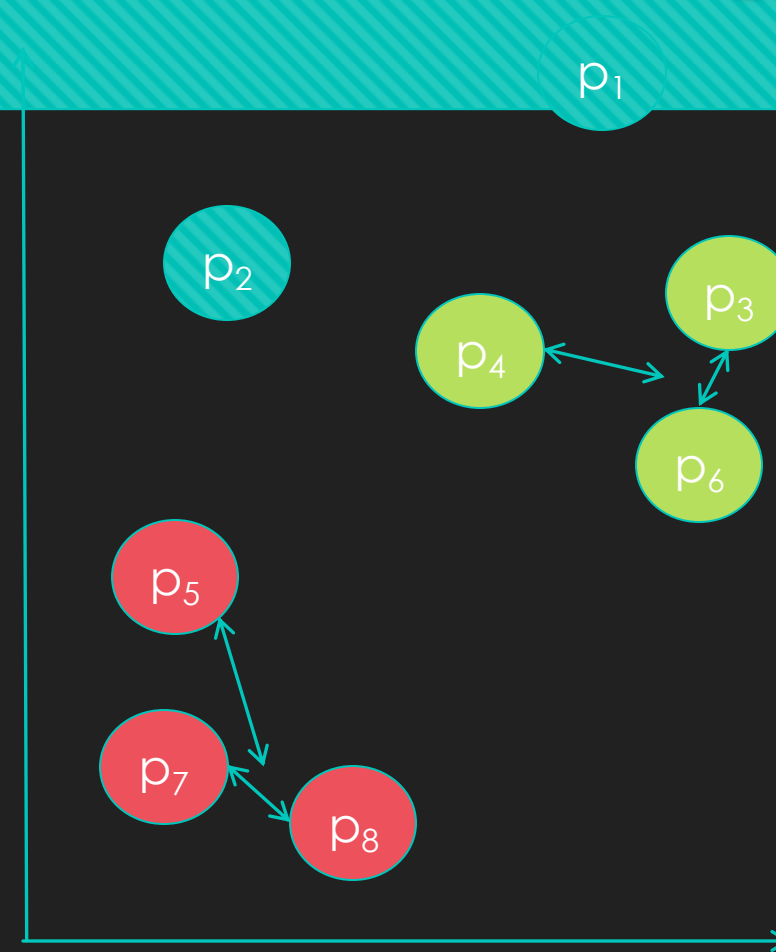
# Agglomerative hierarchical clustering

1. Initialise  $n$  clusters as the  $n$  data points
2. Find closest pair  $\langle p_i, p_j \rangle$  of clusters with distance  $d$
3. while  $d < \text{threshold}$ :
  1. Merge clusters  $\langle p_i, p_j \rangle$
  2. Find closest pair  $\langle p_i, p_j \rangle$  with distance  $d$



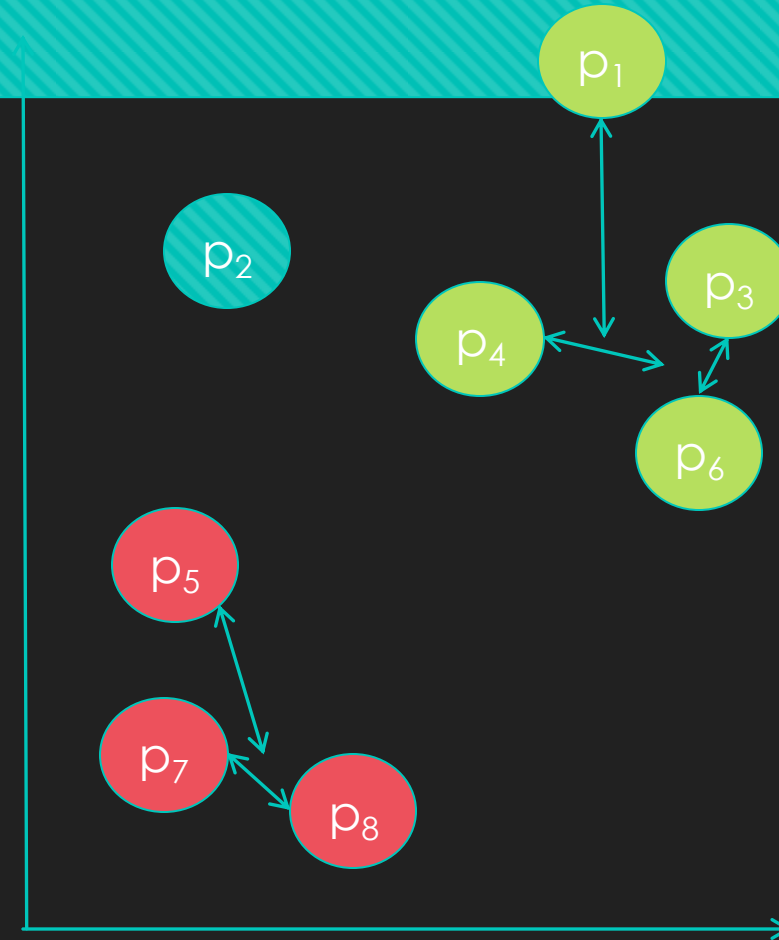
# Agglomerative hierarchical clustering

1. Initialise  $n$  clusters as the  $n$  data points
2. Find closest pair  $\langle p_i, p_j \rangle$  of clusters with distance  $d$
3. while  $d < \text{threshold}$ :
  1. Merge clusters  $\langle p_i, p_j \rangle$
  2. Find closest pair  $\langle p_i, p_j \rangle$  with distance  $d$



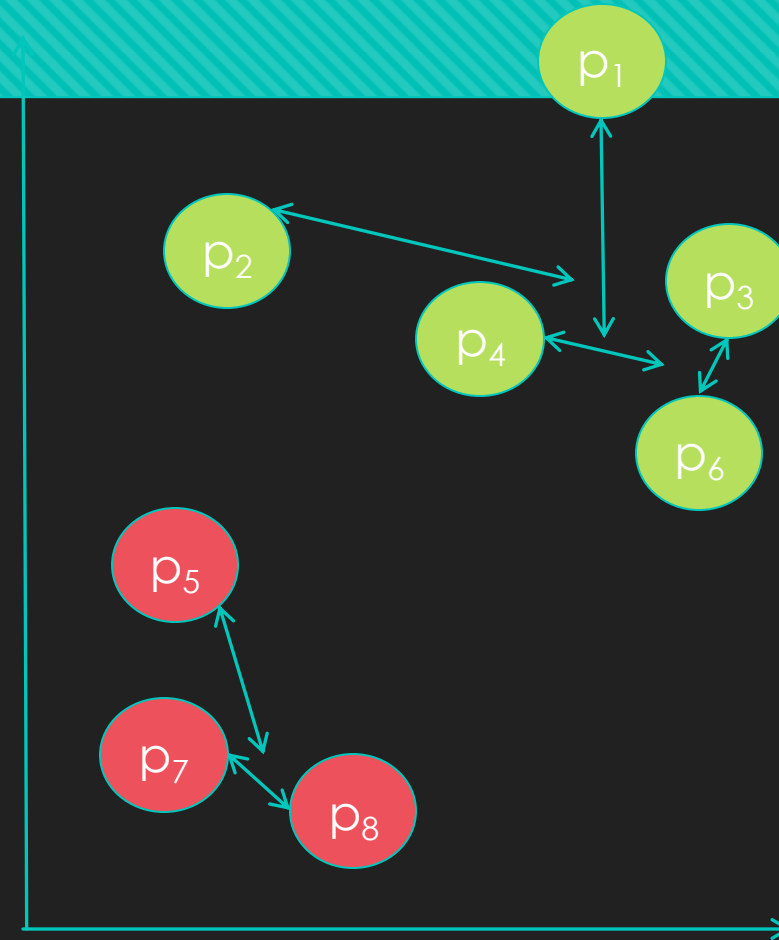
# Agglomerative hierarchical clustering

1. Initialise  $n$  clusters as the  $n$  data points
2. Find closest pair  $\langle p_i, p_j \rangle$  of clusters with distance  $d$
3. while  $d < \text{threshold}$ :
  1. Merge clusters  $\langle p_i, p_j \rangle$
  2. Find closest pair  $\langle p_i, p_j \rangle$  with distance  $d$




# Agglomerative hierarchical clustering

1. Initialise  $n$  clusters as the  $n$  data points
2. Find closest pair  $\langle p_i, p_j \rangle$  of clusters with distance  $d$
3. while  $d < \text{threshold}$ :
  1. Merge clusters  $\langle p_i, p_j \rangle$
  2. Find closest pair  $\langle p_i, p_j \rangle$  with distance  $d$



# Agglomerative hierarchical clustering

1. Initialise  $n$  clusters as the  $n$  data points
2. Find closest pair  $\langle p_i, p_j \rangle$  of clusters with distance  $d$
3. while  $d < \text{threshold}$ :
  1. Merge clusters  $\langle p_i, p_j \rangle$
  2. Find closest pair  $\langle p_i, p_j \rangle$  with distance  $d$





# Disadvantages of agglomerative hierarchical clustering

- Computationally expensive to keep recomputing all nearest neighbours
- runtime =  $O(n^2 \log n)$  where  $n$  is the number of data points
- and the constant is large  $\rightarrow$  multiple of  $d$  where  $d$  is the number of dimensions
- in comparison, k-means is  $O(n * d * k * I)$  where  $k$  is the number of clusters and  $I$  is the number of iterations

# Making progress

- Next week you should complete **all** of the exercises in the single notebook for week 5 on Document Similarity

 Part 1: Lab\_5\_1.ipynb

- And make progress with your coursework