

Week 6: Similarity Analysis

Algorithmic Data Science

2025-26



Warm-up

- Computers store numbers (well everything) in binary
- In decimal, 45 means $4 \times 10^1 + 5 \times 10^0$
- In binary, 1011 means $1 \times 2^3 + 0 \times 2^2 + 1 \times 2^1 + 1 \times 2^0$
- Complete the table below with binary / decimal equivalences

Binary	Decimal
1011	11
101	
1001101	
	32
	100

A byte is 8 bits (where a bit is a 0 or 1). What's the largest number which can be stored in 1 byte?
What's the largest number which can be stored in 4 bytes?

Warm-up

- Computers store numbers (well everything) in binary
- In decimal, 45 means $4 \times 10^1 + 5 \times 10^0$
- In binary, 1011 means $1 \times 2^3 + 0 \times 2^2 + 1 \times 2^1 + 1 \times 2^0$
- Complete the table below with binary / decimal equivalences

Binary	Decimal
1011	11
101	5
1001101	1+4+8+64=77
100000	32
1100100	100

A byte is 8 bits (where a bit is a 0 or 1). What's the largest number which can be stored in 1 byte?
What's the largest number which can be stored in 4 bytes?

Main topics per week

Week	Topic
1	Data structures and data formats
2	Algorithmic complexity. Sorting.
3	Matrices: Manipulation and computation
4	Processes and concurrency
5	Distributed computation
6	Similarity
7	Map/reduce
8	Graphs/networks
9	Graphs/networks, PageRank algorithm
10	Databases
11	<i>independent study</i>

Overview

- applications of similarity / near-neighbour search
- similarity measures
- string similarity
- shingling
- Minhashing
- Locality sensitive hashing (LSH)

Applications of similarity

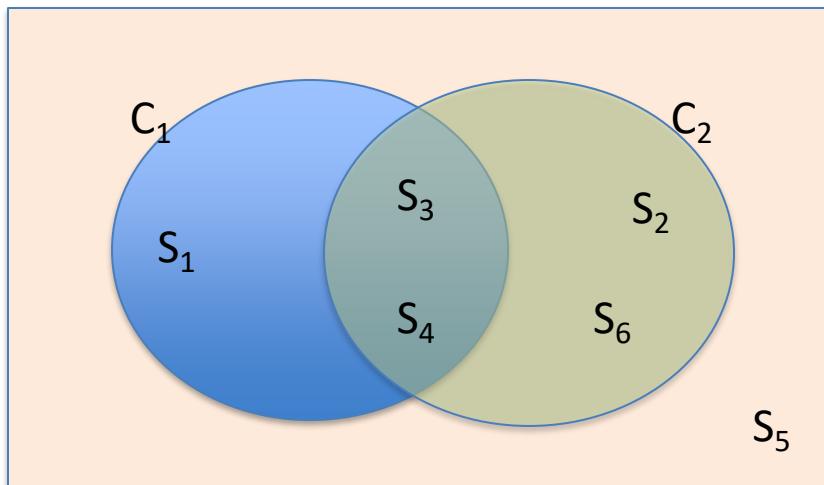
- similarity of documents
 - plagiarism
 - mirror pages
 - articles from the same source
- collaborative filtering
 - online purchases
 - movie ratings
- clustering
 - grouping objects in such a way that objects in the same group (called a **cluster**) are **more similar** to each other than to those in other clusters.

Example

- The ‘objects’ in which we are interested are customers
- We want to consider two customers similar if they have purchased similar items.
- If we have each customer’s purchase history, how do we represent each customer?

Set-theoretic notions of similarity

- Boolean features (a customer C_i either has or hasn't purchased some item S_j) lead naturally to set-theoretic notions of similarity.



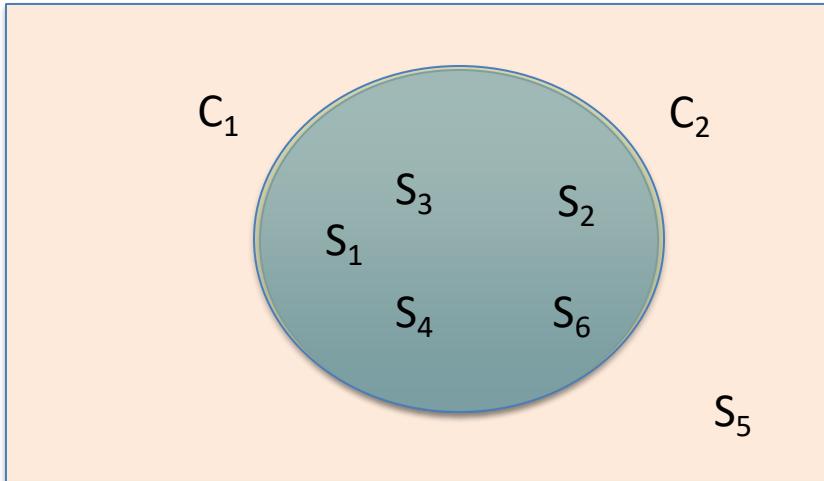
$$C_1 = \{S_1, S_3, S_4\}$$

$$C_2 = \{S_2, S_3, S_4, S_6\}$$

Jaccard's measure is the ratio of the cardinality (size) of the intersection of two sets to the cardinality of the union of two sets

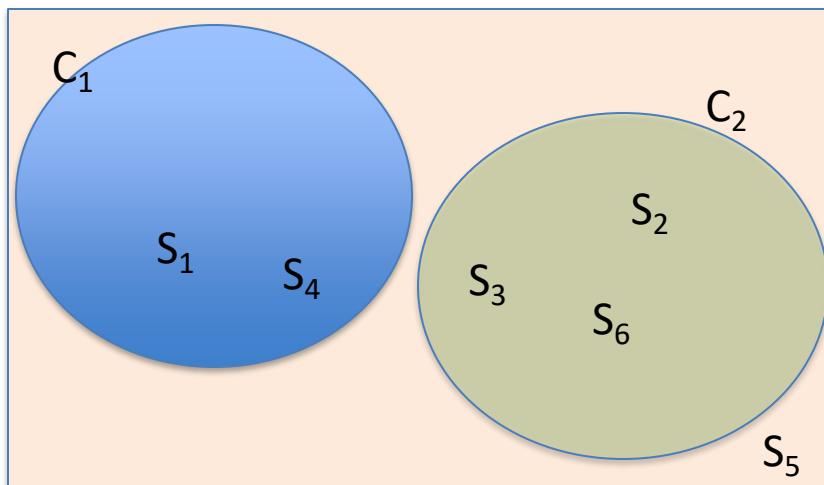
$$Jacc(C_1, C_2) = \frac{|C_1 \cap C_2|}{|C_1 \cup C_2|} = \frac{|C_1 \cap C_2|}{|C_1| + |C_2| - |C_1 \cap C_2|} = \frac{2}{5}$$

Special cases



$$C_1 = C_2 = \{S_1, S_2, S_3, S_4, S_6\}$$

$$\text{Jacc}(C_1, C_2) = 5/5 = 1$$



$$C_1 = \{S_1, S_4\}$$

$$C_2 = \{S_2, S_3, S_6\}$$

$$\text{Jacc}(C_1, C_2) = 0/5 = 0$$

Algorithm for Jaccard's Measure

- What does the run-time for computing Jaccard similarity depend on?
- What are the possible worst-case performances in O notation?

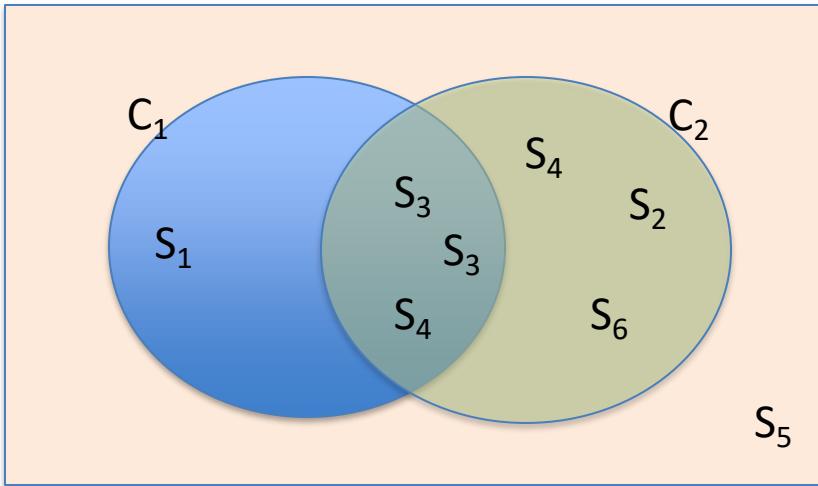
Algorithm for Jaccard's measure

Data stored in Python lists:

```
def jaccard(C1, C2):
    int=0
    union=0
    for item in C1:
        if item in C2:
            int+=1
    union=len(C1)+len(C2)-int
    return int/union
```

Assuming C1 and C2 have length $O(n)$, then this is $O(n^2)$ because the if statement takes $O(n)$ to execute (check every element).

Extending to Bags



$$C_1 = \{S_1, S_3, S_3, S_4\}$$

$$C_2 = \{S_2, S_3, S_3, S_4, S_4, S_6\}$$

Can model duplicate items or even real-valued scores using bags.

The shared (minimum) part of the score goes in the intersection. All of it goes in the union

	C ₁	C ₂
S ₁	1	0
S ₂	0	1
S ₃	2	2
S ₄	1	2
S ₅	0	0
S ₆	0	1

characteristic matrix

$$Jacc(C_1, C_2) = \frac{|C_1 \cap C_2|}{|C_1| + |C_2| - |C_1 \cap C_2|} = \frac{\boxed{\min(S_{i1}, S_{i2})}}{\boxed{i} S_{i1} + S_{i2} - \min(S_{i1}, S_{i2})} = \frac{3}{7}$$

Algorithm for Jaccard's measure

Data stored in a dictionary, bags version of Jaccard:

```
def maketotal(dict1):
    total=0
    for item in dict1:
        total += dict1[item]
    return total

def jaccard(dict1,dict2):
    intersection={}
    for item in dict1.keys():
        if item in dict2.keys():
            intersection[item]=min(dict1[item],dict2[item])

    intersectiontot=maketotal(intersection)
    union = maketotal(dict1)+maketotal(dict2)-intersectiontot
    return intersectiontot/union
```

Assuming dict1 and dict2 have $O(n)$ item, then **if we have no hash collisions:**

the if statement takes $O(1)$ to execute- just look at what is stored at the hash of the item.

Note

- Take care when using Jaccard similarity, to be clear whether you are using the measure applied to **sets** or to **bags**. In general, this choice affects the value you get.

Exercise

- Consider the following 2 sets of items S1 and S2.

$$S1=\{A1, A2, A5, A6\}$$

$$S2=\{A2, A3, A5\}$$

What is the Jaccard similarity of sets S1 and S2?

- (a) 3/4
- (b) 2/5
- (c) 2/7
- (d) 1/3

Exercise (Solution)

- Consider the following 2 sets of items S1 and S2.

$$S1 = \{A1, A2, A5, A6\}$$

$$S2 = \{A2, A3, A5\}$$

Compute the Jaccard similarity of sets S1 and S2.

Intersection = {A2, A5} size of intersection = 2

Union = {A1, A2, A3, A5, A6} size of union = 5

Jaccard similarity = $|I| / |U| = 2/5 = 0.4$

Compute the cosine similarity of sets S1 and S2.

	S1	S2	$S1 \cdot S2 = 0 + 1 + 0 + 1 + 0 = 2$
A1	1	0	$S1 \cdot S1 = 4$
A2	1	1	$S2 \cdot S2 = 3$
A3	0	1	
A5	1	1	
A6	1	0	

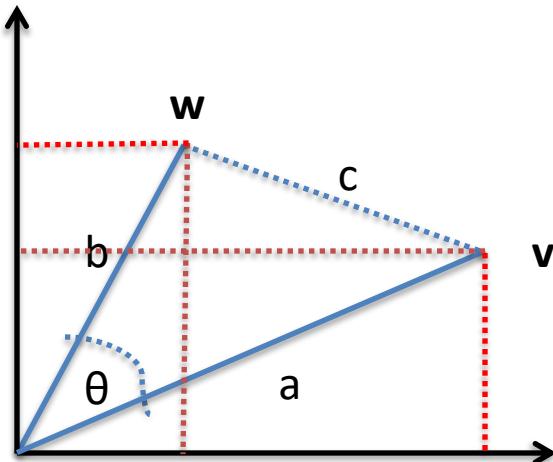
$$\text{Cos} = \frac{S1 \cdot S2}{\sqrt{(S1 \cdot S1)(S2 \cdot S2)}} = \frac{2}{\sqrt{12}} = \frac{2}{2\sqrt{3}} = \frac{1}{\sqrt{3}} \approx 0.58$$

Cosine similarity

This makes use of the **dot product** for vectors:

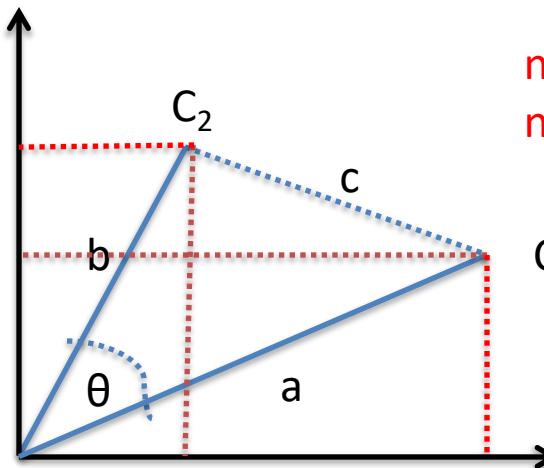
$$\mathbf{v} \cdot \mathbf{w} = \sum_{i=1}^n v_i w_i$$

$$\cos(\theta) = \frac{\mathbf{v} \cdot \mathbf{w}}{\sqrt{(\mathbf{v} \cdot \mathbf{v})(\mathbf{w} \cdot \mathbf{w})}}$$



Cosine similarity

- Real valued ‘vector’ representations of objects lead naturally to geometric notions of similarity



n dimensions,
n = 6 here

$$\begin{bmatrix} & C_1 & C_2 \\ S_1 & 1 & 0 \\ S_2 & 0 & 1 \\ S_3 & 2 & 2 \\ S_4 & 3 & 1 \\ S_5 & 1 & 1.5 \\ S_6 & 0 & 1 \end{bmatrix} \xrightarrow{C_1 \cdot C_2} \begin{array}{|c|} \hline 0 \\ 0 \\ 4 \\ 3 \\ 1.5 \\ 0 \\ \hline \Sigma = 8.5 \end{array}$$

$$\cos(C_1, C_2) = \frac{C_1 \cdot C_2}{\sqrt{C_1 \cdot C_1 \square C_2 \cdot C_2}}$$

$C_1 \cdot C_2$ is the dot product.
Also known as the inner product $\langle C_1, C_2 \rangle$ or the scalar product

Algorithm for Cosine Measure

- What does the running time of this algorithm depend on?
- Give an estimate of its worst-case performance in O notation

```
def naiveCosine (a , b):  
    num=0  
    d1=0  
    d2=0  
    for i in range len( a ) :  
        num += a [ i ] *b [ i ]  
        d1 += a [ i ] *a [ i ]  
        d2 += b [ i ] *b [ i ]  
    return num / ( d1*d2 ) **0.5
```

Finding similarity of two strings

- Strings can be modelled as bags-of-characters
- Long strings (documents!) can be modelled as bags-of-words

“colour” -> {c:1, o:2, l:1, u:1, r:1}
“color” -> {c:1, o:2, l:1, r:1}

a dictionary is a compact sparse representation for a bag but it is equivalent to the dense matrix representation

Exercise (at home): Show that $\text{Jacc}(\text{“colour”, “color”}) = 5/6$

Question:
 $\cos(\text{“colour”, “color”}) = ?$

$$(a) \sqrt{\frac{8}{7}}$$

$$(c) \frac{8}{7} \quad (d) \frac{7}{8}$$

	S1	S2
C	1	1
O	2	2
L	1	1
U	1	0
R	1	1

$$S1 \cdot S2 = 1+4+1+1=7$$

$$S1 \cdot S1 = 1+4+1+1+1=8$$

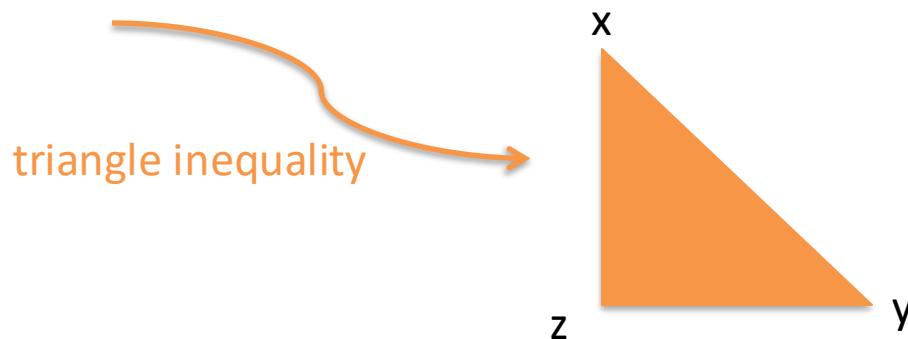
$$S2 \cdot S2 = 1+4+1+1=7$$

$$\cos = \frac{S1 \cdot S2}{\sqrt{(S1 \cdot S1)(S2 \cdot S2)}} = \frac{7}{\sqrt{7 * 8}} = \sqrt{\frac{7}{8}}$$

Similarity vs Distance

- Distance measures measure dissimilarity

distance measures	similarity measures
$d(x,y) \geq 0$	$0 \leq \text{sim}(x,y) \leq 1$
$d(x,y) = 0 \text{ iff } x=y$	$\text{sim}(x,y)=1 \text{ iff } x=y$
$d(x,y) = d(y,x)$	$\text{sim}(x,y) = \text{sim}(y,x)$
$d(x,y) \leq d(x,z) + d(z,y)$	



Correlation vs Cosine Similarity

To compute correlation of 2 variables X and Y

- Subtract the mean of X from each value X_i and the mean of Y from each value Y_i
- Compute the dot-product of the transformed X and Y. This is the **covariance** of X and Y
- Divide by the square root of the product of $\text{cov}(X,X)$ and $\text{cov}(Y,Y)$

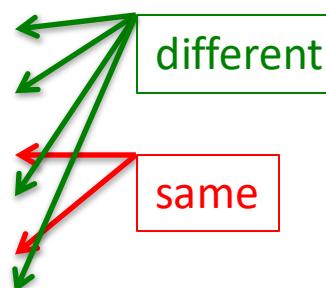
To compute cosine similarity between 2 vectors X and Y

- Compute the dot product of X and Y : $\langle X, Y \rangle$
- Divide by the square root of the product of $\langle X, X \rangle$ and $\langle Y, Y \rangle$

The only difference is that when computing correlation, we compute covariance rather than a simple dot product i.e., we standardize by subtracting the means first.

Other measures: Hamming distance

- The number of vector components (dimensions) in which two objects differ.
- Usually only applied to Boolean vectors (e.g., sets) but can be applied to bags

	C_1	C_2	
S_1	1	0	
S_2	0	1	
S_3	2	2	
S_4	1	2	
S_5	0	0	
S_6	0	1	

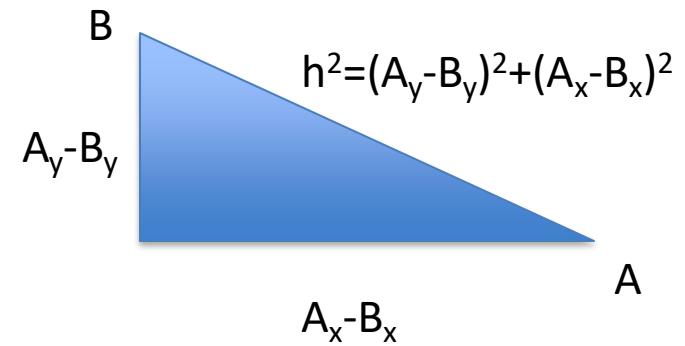
Hamming distance(C_1, C_2) = 4

L Norms

Most people are familiar with the L_2 Norm (also known as the **Euclidean distance**), which is Pythagoras theorem in n-dimensions:

$$L_2(A, B) = \sqrt{\sum_{i=1}^n (A_i - B_i)^2}$$

In general, the L_k Norm is given by



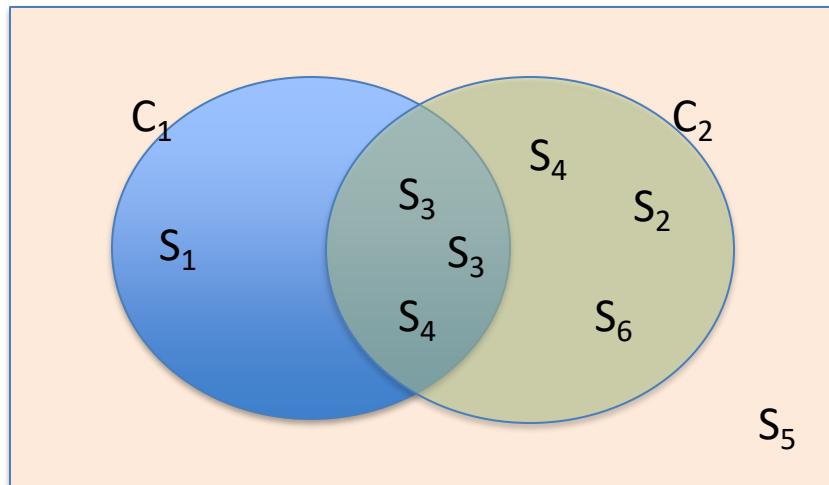
$$L_k(A, B) = \sqrt[k]{\sum_{i=1}^n |A_i - B_i|^k}$$

The L_1 Norm (or Manhattan or City Block) distance is

$$L_1(A, B) = \sum_{i=1}^n |A_i - B_i|$$

Probabilistic measures of similarity

- Frequencies can be easily converted into probabilities



What is the probability that a randomly chosen item is S_i given it is in the set / bag C_j ?

	C_1	C_2		C_1	C_2	
S_1	1	0	S_1	0.25	0	
S_2	0	1	S_2	0	0.166	
S_3	2	2	\rightarrow	S_3	0.5	0.333
S_4	1	2	S_4	0.25	0.333	
S_5	0	0	S_5	0	0	
S_6	0	1	S_6	0	0.166	

Probabilistic measures of similarity

Most well-known ‘distance’ measure for probability distributions is the Kullback-Leibler divergence measure

$$D_{KL}(C_1 \parallel C_2) = \sum_i p_{i1} \log \frac{p_{i1}}{p_{i2}}$$

What is the average penalty (i.e., difference in log probabilities) if you use the distribution for C_1 in place of the distribution for C_2 ?

This is not strictly a distance measure because it is not symmetric. The Jenson-Shannon divergence measure is the symmetric version, which measures distance as the average Kullback-Leibler divergence to the centroid of the distributions.

$$JS(A, B) = \frac{1}{2} (KL(A, M) + KL(B, M))$$

$$\text{where } M = \frac{1}{2}(A + B)$$

Disadvantages of Using Bags for Text

- Not sensitive to order
 - “brag” = “grab”
- If applied to documents where the atomic units are words
 - does not capture relationships between different words

The old man chased the small dog that bit a naughty child.



The old dog chased the naughty small child that bit a man.

Shingling

The old man chased the
small dog that bit a
naughty child.



The old dog chased the
naughty small child that
bit a man.

- A bag-of-words (or bag of characters) representation will lead to these strings being considered identical.
- We could use a bag-of-*n*grams :
 - unigram = 1 word, bigram = 2 words, trigram = 3 words, ngram = n words
 - ***Exercise:*** Show that the Jaccard similarity if we used a bag of bigrams is $2/9$.
- Alternative is to use a set or bag of shingles. A k -shingle for a document is any string of length k found in the document

Shingling

Example: What are the sets of 3-shingles for the strings “john loves mary” and “mary loves john”?

If a string has m characters, it will have at most m-k distinct shingles

A: john loves mary			B: mary loves john		
joh	ohn	hn_	mar	ary	ry_
n_!	_lo	lov	y_!	_lo	lov
ove	ves	es_	ove	ves	es_
s_m	_ma	mar	s_j	_jo	joh
ary			ohn		

$$\text{Jacc}(A,B) = 9/(13+13-9) = 9/17$$

Edit distance

- The edit distance between strings X and Y is the smallest number of operations required to transform X into Y, where the operations allowed are insertion and deletion (and also sometimes transposition and mutation).
- There are variants where the different operations have different costs but lets assume cost of each operation = 1

X	Y		$d(X,Y)$
colour	color	delete c_5	1
doggy	daddy	delete c_2 , delete c_3 , delete c_4 , insert 'a' at c_2 , insert 'd' at c_3 , insert 'd' at c_4	6
brag	grab	??	4
house	home	??	?

Similarity for large collection of documents

documents

	C1	C2	C3	C4
S1	1	0	0	1			
S2	0	0	1	0			
S3	0	1	0	1			
S4	1	0	1	1			
S5	0	0	1	0			
...							
...							
...							

- We're going to see an efficient way of analysing similarity for all pairs of documents in a collection.
- It will use the **set version of Jaccard similarity**.
 - So want to choose a shingle length where most shingles don't occur in most documents, so very different documents have very different shingle sets.

Choosing the shingle size

- k should be picked large enough that the probability of any given shingle appearing in any given document is low
- depends on the length of the typical document and the size of the character set

Example: if our corpus of documents is emails then $k=5$ is probably appropriate. Why?

- Assume number of characters is 27.
- Number of shingles = $27^5 = 14, 348, 907$
- Typical email length << 14, 348 907 characters ☺
- In practice, there are more than 27 characters but many are very rare which increases probability of shingles of more common letters
- So actually better to assume number of characters for English ≈ 20

Shingles vs Bags-of-words

Representation	parameters	dimensionality
shingles	characters = 27 k = 5	$27^5 \approx 1.4 \times 10^7$
shingles	characters = 20 k = 9	$20^9 = 5.12 \times 10^{11}$
bag-of-words	vocabulary = 500K	5×10^5
bag-of-ngrams	vocabulary = 500K n = 2	$500,000^2 = 2.5 \times 10^{11}$

- Shingles are fixed length whereas words are variable in length

Hashing shingles

- The ASCII character set has 128 characters
- If we use 1 byte per character, a 9-shingle will take 9 bytes
- And many of the possible shingles will never occur
- Use a hash function which maps 9-shingles to numbers in range $0 \rightarrow 2^{32} - 1$
- Then likelihood of each hashed value occurring much more equal.
- And such a number can be stored in 4 bytes
- However, still have several times more shingles per document as individual characters – need compression if lots of documents to analyse.

Minhashing

- A technique for constructing small **signatures** from large sets whilst preserving estimates of similarity.

Algorithm for minhashing a set represented by a column of a characteristic matrix:

1. pick a permutation of the rows
2. The minhash value of any column is the first row in the permuted order in which the column has a 1.
3. Repeat m times to get a minhash signature of length m

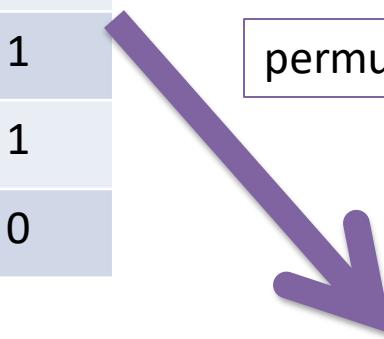
	C1	C2	C3	C4
S1	1	0	0	1
S2	0	0	1	0
S3	0	1	0	1
S4	1	0	1	1
S5	0	0	1	0

permutation 1



	C1	C2	C3	C4
S2	0	0	1	0
S5	0	0	1	0
S3	0	1	0	1
S1	1	0	0	1
S4	1	0	1	1
MH	4	3	1	3

permutation 2



	C1	C2	C3	C4
S4	1	0	1	1
S3	0	1	0	1
S2	0	0	1	0
S1	1	0	0	1
S5	0	0	1	0
MH	1	2	1	1

How many permutations?

- In practice, m will be much less than the dimensionality of the matrix (and much much less than the number of possible permutations)
- How large should m be? Say $m = 100$
- By minhashing we have reduced the dimensionality of the characteristic matrix :
- Originally: e.g. 2^{32} Boolean values, each Boolean takes 1 bit so 2^{24} bytes (=17MB) per column
- In minhash signature: each integer $< 2^{32}$ so can be stored in 4 bytes so 400 bytes per column

Computing Minhash Signatures

- Not feasible to permute a large matrix explicitly
 - would have to pick a random permutation of billions of rows
 - then sort all of those rows ...
- Simulate the effect of a random permutation using a hash function, $h(r)$
- Same number of buckets as rows
- Whilst there will be some collisions, we can maintain the fiction that our hash function h permutes row r to position $h(r)$
- So instead of m random permutations, randomly choose m hash functions on the rows

Computing Minhash signatures

1. LET $M_{r,c}$ be element of the characteristic matrix for the r th element for c th set.
2. Let $SIG_{i,c}$ be the element of the signature matrix for the i th hash function and column c .
3. Initialise $SIG_{i,c}$ to ∞ for all i and c
4. FOR each row r :
 FOR each hash function h_i :
 compute $h_i(r)$
 FOR each column c :
 IF c has a 0 in row r : do nothing
 ELSE: $SIG_{i,c} = \text{MIN}(SIG_{i,c}, h_i(r))$

Minhashing and Jaccard Similarity

The probability that the minhash function for a random permutation of rows produces the same value for two sets equals the Jaccard similarity of those sets.

M	C1	C2	C3	C4
S1	1	0	0	1
S2	0	0	1	0
S3	0	1	0	1
S4	1	0	1	1
S5	0	0	1	0



SIG	C1	C2	C3	C4
MH1	4	3	1	3
MH2	1	2	1	1

Why?

Minhashing and Jaccard

	C1	C2	C3	C4
S1	1	0	0	1
S2	0	0	1	0
S3	0	1	0	1
S4	1	0	1	1
S5	0	0	1	0

For any given pair of columns:

- Type X rows have a 1 in both
- Type Y rows have different values
- Type Z rows have a 0 in both

For sparse matrices, most rows for most pairings will be type Z

It is the ratio of type X to type Y rows that determine $Jacc(C_i, C_j)$ and also the probability that $h(C_i) = h(C_j)$

$$Jacc(C_3, C_4) = \frac{X_{3,4}}{X_{3,4} + Y_{3,4}} = \frac{1}{5}$$

	C1	C2	C3	C4
S4	1	0	1	1
S3	0	1	0	1
S2	0	0	1	0
S1	1	0	0	1
S5	0	0	1	0
MH	1	2	1	1

In a random permutation, the probability that we meet a type X row before we meet a type Y row is also $X_{3,4}/(X_{3,4}+Y_{3,4})$

If we do meet a type X row before we meet a type Y row, then we get $MH(C_3) = MH(C_4)$

	C1	C2	C3	C4
S2	0	0	1	0
S5	0	0	1	0
S3	0	1	0	1
S1	1	0	0	1
S4	1	0	1	1
MH	4	3	1	3

However, if we meet a type Y row before we meet a type X row then we get $MH(C_3) \neq MH(C_4)$

Minhashing and Jaccard Similarity

SIG	C1	C2	C3	C4
MH1	4	3	1	3
MH2	1	2	1	1

So, if we carried out ALL random permutations, the proportion of matches in the minhash signatures for 2 objects would equal Jaccard similarity

For a random selection of permutations, proportion of matches will estimate the Jaccard similarity

estimate of Jaccard	C1	C2	C3	C4
C1	1	0	0.5	0.5
C2	0	1	0	0.5
C3	0.5	0	1	0.5
C4	0.5	0.5	0.5	1

Efficient similarity analysis: Summary

1. Construct a characteristic matrix:
 - columns are population members e.g. documents, customers
 - rows are items, e.g. hashed k-shingles, items for sale
2. Compute m minhash signatures.

Jaccard similarity of C_i and C_j is approximately the proportion of minhash signatures which agree in column i and column j

But if n is the size of the population (number of documents / customers etc), it is still $O(n^2)$ to do all pairs-similarity.

Locality-Sensitive Hashing (LSH)

- A technique for efficiently finding nearest neighbours without computing all-pairs similarities.
- General approach is to hash items several times in such a way that similar items are more likely to be hashed to the same bucket than dissimilar items
- Any pair hashed to the same bucket for any of the hashings is then considered a **candidate pair**
- Only compute similarities for candidate pairs.
- There will be false positives but these will be found whilst computing similarities
- There will be false negatives (candidate pairs completely missed) – need to minimise these as no way of recovering them.

LSH for a Minhash Signature

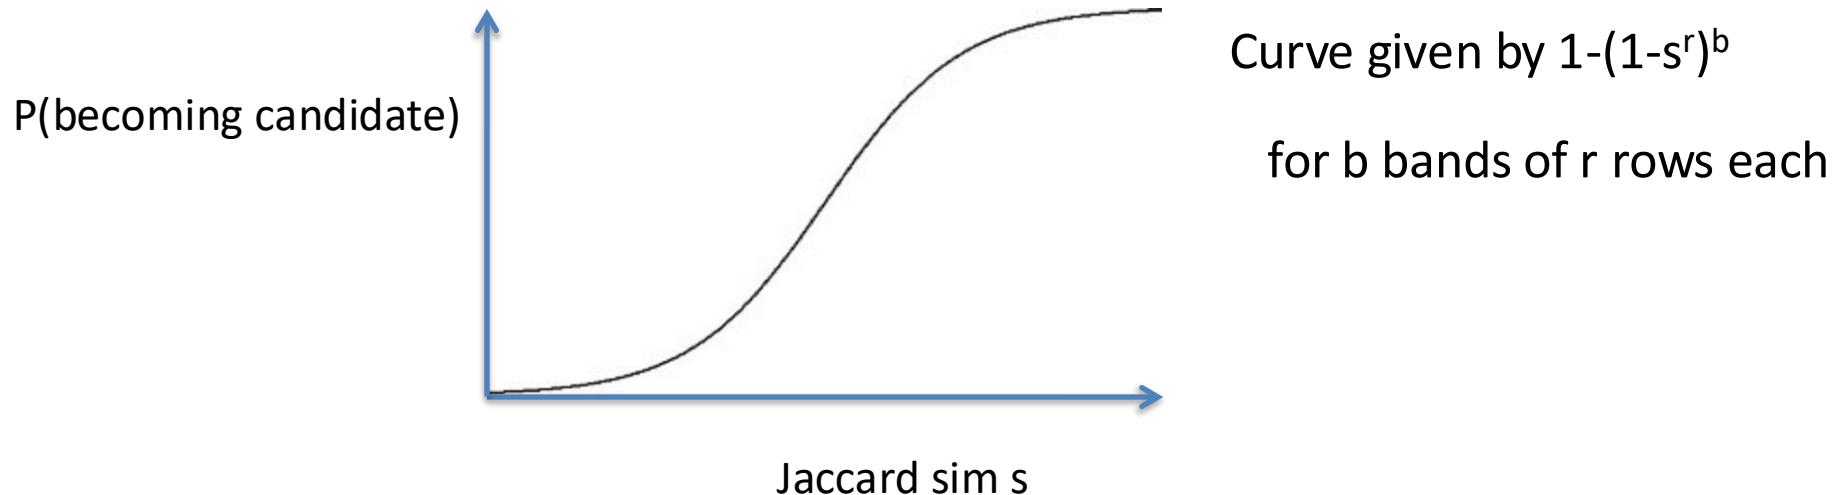
	C1	C2	C3	C4	C5
band 1	4	3	1	3	1
	1	2	1	1	1
	0	1	3	1	3
band 2					
band 3					
band 4					

- Each hash function only considers a **band** of rows.
- Columns which are identical in the rows of a particular band must be hashed to the same bucket for that band
- There will be accidental collisions (leading to false positives)
- However, similar items will probably be identical in at least 1 band

Analysis of Banding Technique

- Suppose we use b bands of r rows each and that a particular pair of documents have Jaccard similarity s
- $P(\text{sig} \text{ agree in all rows of a particular band}) = s^r$
- $P(\text{sig} \text{ do not agree in all rows of a particular band}) = 1 - s^r$
- $P(\text{sig} \text{ do not agree in all rows of any band}) = (1 - s^r)^b$
- $P(\text{sig} \text{ agree in all rows of at least one band}) = 1 - (1 - s^r)^b$

Analysis of Banding Technique



- The threshold Jaccard similarity at which it becomes likely that the pair will become a candidate depends on b and r .
- The more rows per band, the higher this threshold is.
- An approximation to the threshold (where $\text{Prob}=1/2$) is threshold similarity = $(1/b)^{1/r}$
- If 100 rows are divided into 20 bands of 5, what will the threshold similarity be?
- What if we use 5 bands of 20?

Complexity of LSH

- Assuming that parameters have been chosen so that only 10% of pairs are considered to be candidate pairs by LSH, how much efficiency saving do you get for finding the k-nearest neighbours?

If n is now number of documents,

- $O(n)$ to minhash.
- $O(n)$ to do LSH (apply hash to each of n documents, and then look at what is in each bucket).
- So, still $O(n^2)$, but with factor of 10 saving on the constant

Efficient similarity analysis: Summary

1. Construct a characteristic matrix:
 - columns are population members e.g. documents, customers
 - rows are items, e.g. hashed k-shingles, items for sale
2. Compute m minhash signatures.

Jaccard similarity of C_i and C_j is approximately the proportion of minhash signatures which agree in column i and column j

1. Choose a similarity threshold, t
2. Construct candidate pairs by applying LSH
3. Compute similarities for candidate pairs from minhash signatures
4. Check similarity for a few original documents (to verify nothing went wrong).

Bonus material: Other Similarity Measures and LSH

- No guarantee that a particular distance / similarity measure has a locality-sensitive family of hash functions
- However possible to do so for:-
 - Hamming distance
 - Cosine distance
 - Euclidean distance

Bonus material: LSH for Cosine

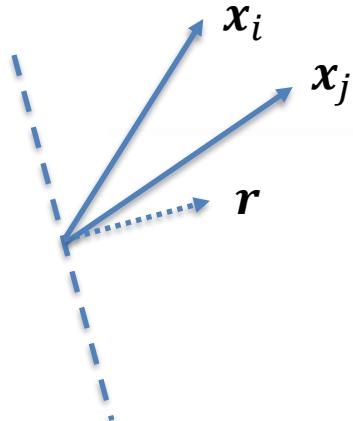
- Consider two points (described by position vectors) and a random hyperplane through the Origin
- The two points are either on the same side of the hyperplane or on different sides of the hyperplane
- If we take the dot product of each vector with the normal vector to the plane
 - Same sign -> same side of hyperplane
 - Different signs -> different sides of hyperplane

Bonus material: LSH for Cosine

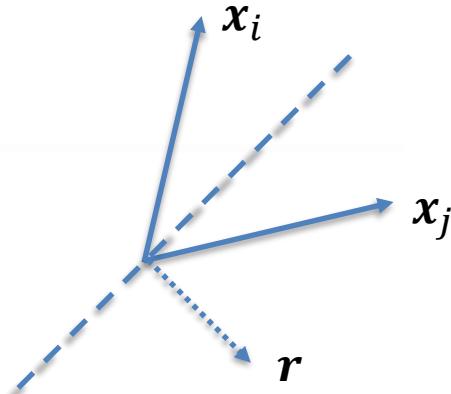
- The probability that a *random hyperplane* separates two unit vectors depends on the angle between them:

$$\Pr[\text{sign}(\mathbf{x}_i^T \mathbf{r}) = \text{sign}(\mathbf{x}_j^T \mathbf{r})] = 1 - \frac{1}{\pi} \cos^{-1}(\mathbf{x}_i^T \mathbf{x}_j)$$

High dot product:
unlikely to split



Lower dot product:
likely to split



Corresponding hash function:

$$h_{\mathbf{r}}(\mathbf{x}) = \begin{cases} 1, & \text{if } \mathbf{r}^T \mathbf{x} \geq 0 \\ 0, & \text{otherwise} \end{cases}$$

What have you learnt about the following topics?

- applications of near-neighbour search
- similarity and distance measures
- string similarity
- shingling
- min-hashing
- LSH